

Akshat Parikh

San Jose, CA | (408) 637-9861 | akshat.sparikh@gmail.com | [linkedin.com/in/parikhakshat](https://www.linkedin.com/in/parikhakshat) | github.com/AkshatParikh16

Education

Master's in Data Science, The University of Texas at Arlington - **GPA: 4.00/4.00** **Aug 2024 - May 2026**
Relevant Courses: Machine Learning, AI, Deep Learning, Project Management Texas, United States

Bachelor's in Information Technology, Charusat University - **GPA: 3.98/4.00** **Sept 2020 - May 2024**
Relevant Courses: Big Data Analytics, Statistical Modeling, Data Mining, Cloud Computing Gujarat, India

Skills

Programming: Python, SQL, R

Frameworks & Libraries: Scikit-learn, Pandas, NumPy, PyTorch, TensorFlow, Keras, FastAPI, PySpark

GenAI: LangChain, LangGraph, HuggingFace, Pinecone, RAG, LLMs, Agentic Workflow, Prompt Engineering, Langfuse

Data Science: Model Evaluation & Optimization, Data Modeling, Experimental Design, A/B Testing, Fine-tuning, Feature Engineering, Predictive Modeling, Model Deployment, Hypothesis Testing

Visualization: Power BI, Tableau, Excel, Matplotlib, Seaborn, Plotly

Other Tools: Git, GitHub Actions, CI/CD, Docker, AWS, Databricks

Professional Experience

Data Scientist (Student Assistant) **May 2025 - Apr 2026**
UTA Career Development Center (Employer Outreach Services) Arlington, Texas

- Automated ETL pipelines centralizing 500K+ employer, student, and application records from Handshake and HubSpot CRM into a PostgreSQL data model on AWS S3 and RDS, reducing manual reporting by 35%.
- Developed an employer intelligence platform on AWS EC2 to track partner engagement, application activity, and outreach performance metrics, eliminating ad-hoc data requests and enabling real-time insights.
- Integrated an AI Assistant into the platform, letting staff retrieve reports, cutting inquiry response time by 40%.
- Partnered with program teams to analyze outreach trends and redesign campaign, increasing event participation by 12%.

Data Science Lead **Jan 2026 - Apr 2026**
UTA Office of Information Technology (Capstone Project) Arlington, Texas

- Led a 3-member team to deliver a Power BI reporting system for UTA's Technology Department, unifying fragmented manual Excel files and PDFs into a single real-time dashboard, saving 80+ hours of manual data management per month.
- Streamlined lab performance reporting across 11 campus labs and presented findings biweekly, giving OIT Managers and Associates real-time visibility into operations.

Machine Learning Engineer Intern **Sept 2023 - May 2024**
Agevole Innovation Gujarat, India

- Contributed to an NLP-based candidate matching system, developing a text preprocessing pipeline using spaCy and NLTK for tokenization, lemmatization, and NER, enhancing candidate-to-role matching accuracy.
- Evaluated retrieval approaches by comparing SBERT and TF-IDF across 10K+ resumes on precision, latency, and memory tradeoffs, documenting findings and recommendations accepted by the engineering team.
- Implemented Flask API endpoints to fetch live job postings from external sources, enabling candidates to receive real-time role recommendations aligned with their skills and experience.
- Identified patterns to reduce early-stage candidate drop-off through cross-functional collaboration with recruiter teams.

Hackathon Projects

Agentic Landing Page Builder | [GitHub](#) **WeaveHacks 4**

- Architected an agentic system by designing a CEO agent in LangGraph that orchestrates 5 specialized sub-agents across OpenAI models, each evaluated against output quality metrics to iteratively refine landing page generation.
- Incorporated Weights & Biases Weave, providing full observability across token usage, agent outputs, and execution traces.

Projects

InspectAI | **LangGraph, Prometheus, Claude, OpenAI, FastAPI, DuckDB, LiteLLM, Streamlit** | [GitHub](#)

- Built a production-ready testing framework that generates scenarios against AI systems and classifies failures using LLM-as-judge scoring, cutting manual QA review time from hours to minutes and surfacing fix recommendations via dashboard.

DocChat | **Python, LLaMA3, LangChain, HuggingFace, Pinecone, FastAPI** | [GitHub](#)

- Engineered a RAG-based AI chatbot with built-in accuracy checks that answers user questions from uploaded documents.

Customer Churn Prediction | **Python, LightGBM, CatBoost, XGBoost, Docker, AWS (ECR, EC2)** | [GitHub](#)

- Trained and tracked predictive churn models in MLflow, achieving 0.85 ROC-AUC on 51K customer records.
- Deployed the model as a production API on AWS using Docker and GitHub Actions, automating the release process.