

Akshat Parikh

San Jose, CA | (408) 637-9861 | akshat.sparikh@gmail.com | [linkedin.com/in/parikhakshat](https://www.linkedin.com/in/parikhakshat) | [Portfolio](#)

EDUCATION

Master's in Data Science, The University of Texas at Arlington - **GPA: 4.00/4.00**

Relevant Courses: Machine Learning, Generative AI, System Designs, Deep Learning

Aug 2024 - May 2026

Texas, United States

Bachelor's in Information Technology, Charusat University - **GPA: 3.98/4.00**

Relevant Courses: Statistical Modeling, Data Structures, Database Management, Cloud Computing

Sept 2020 - May 2024

Gujarat, India

SKILLS

Core Competencies: AI Architectures, Backend Development, Distributed Systems, Cloud Infrastructure, Algorithm Design, LLMs, RAG, Agentic Workflows, Prompt Engineering, MCP, Fine-tuning, vLLM, Inference Serving, Model Evaluation & Optimization, Integration Testing, Experimental Design, A/B Testing, Predictive Modeling, MLOps, Feature Engineering, Data Pipelines, Data Mining, Data Visualization, NLP, Shell Scripting, Operating Systems (UNIX/Linux), APIs, CI/CD

Languages & Frameworks: Python, SQL, R, C++, Go, PowerShell, Java, Scala, C#, JavaScript, HTML/CSS, LangChain, LangGraph, LlamaIndex, CrewAI, AutoGen, Agent SDK, Hugging Face, PyTorch, TensorFlow, Keras, Scikit-learn, Pandas, NumPy, FastAPI, PySpark, Kafka, JAX, Django, Flask, Spark, React, Matplotlib, Plotly, Redux, jQuery, Pytest

Tools/Platforms: AWS, GCP, Azure, Docker, Kubernetes, Git (Version Control), MLflow, Databricks, Vector Databases (Pinecone, ChromaDB), PostgreSQL, MySQL, Redis, Tableau, Power BI, Terraform, Snowflake, BigQuery, Hadoop, Postman

PROFESSIONAL EXPERIENCE

Graduate Data Scientist (Employer Outreach Services)

May 2025 - Apr 2026

UTA Career Development Center

Arlington, Texas

- Automated ETL pipelines consolidating **500K+** employer, student, and application records from Handshake and HubSpot CRM into a single **PostgreSQL** data model on **AWS S3** and **RDS**, **cutting manual reporting time by 35%**.
- Developed an **employer intelligence platform** in **Power BI**, deployed on **AWS EC2**, to track partner engagement, application activity, and outreach performance in one place, **replacing ad-hoc data requests with real-time insights**.
- Integrated an **AI Assistant** into the platform, cutting staff work down to a simple data upload every **2-month cycle** instead of manual compilation, **reducing recurring staff workload by 50%**.
- Proposed a new outreach channel, podcast, and redesigned campaign messaging, then **collaborated cross-functionally** with program teams to execute it, **driving a 25% increase in event participation**.

Student Analytics Engineer

Jan 2026 - Apr 2026

UTA Office of Information Technology

Arlington, Texas

- Led a **3-member team** to build a **Power BI** reporting system that replaced fragmented Excel files and PDF reports with one live dashboard, saving the department **80+ hours of manual data management a month**.
- Tracked lab performance metrics across UTA's technology labs and presented findings **biweekly** to OIT Managers and Associates, giving **leadership a consistent read on where issues were emerging**.

Machine Learning Engineer Intern

Sept 2023 - May 2024

Agevole Innovation

Gujarat, India

- Contributed to an **NLP-based candidate matching system**, developing a text preprocessing pipeline using **spaCy** and **NLTK** for tokenization, lemmatization, and NER, **enhancing candidate-to-role matching accuracy**.
- Benchmarked **SBERT** against **TF-IDF** across **10K+** resumes on precision, latency, and memory tradeoffs, and wrote up what actually held up, getting the findings **adopted by the engineering team**.
- Implemented **Flask API endpoints** pulling in live job postings from external sources, delivering **skill-aligned role recommendations for candidates**.

PROJECTS

InspectAI | LangGraph, Prometheus, Claude, OpenAI, FastAPI, DuckDB, LiteLLM, Streamlit | [GitHub](#)

- Engineered a **production-ready testing framework** using **50 dynamic personas** to generate scenarios for **RAG and multi-agent systems**, scored by an unbiased **3-judge LLM panel** to reduce single-model bias.
- Clustered outputs into distinct **failure modes**, generating **guardrails** and **cutting QA review time from hours to minutes**.

E-Commerce Analytics & Pricing Platform | FastAPI, Kafka, PySpark, ClickHouse, dbt, Airflow, Docker

- Architected a **real-time, fault-tolerant** data pipeline pairing **FastAPI** with **Kafka** to process live clickstream events at **sub-second latency** across a **distributed, containerized streaming layer** for **cloud-native deployment**.
- Built a pricing engine with **PySpark Streaming** and **ClickHouse**, orchestrating **scalable batch workflows** in **dbt** and **Airflow**.

DocChat | Python, LangChain, HuggingFace, Pinecone, LLaMA3, Ollama, FastAPI | [GitHub](#)

- Constructed a **RAG-based document assistant** that pairs **Pinecone** vector search with **BM25** keyword retrieval, backed by a **built-in confidence check**, to answer questions straight from uploaded documents.

Customer Churn Prediction | Python, LightGBM, CatBoost, XGBoost, MLflow, Docker, AWS | [GitHub](#)

- Trained **predictive churn models** on **51K customer records**, benchmarking multiple algorithms to reach **0.85 ROC-AUC**.
- Logged experiments in **MLflow** and deployed the model as a production API on **AWS EC2** with **Docker** and **CI/CD**.

Agentic Landing Page Builder (Hackathon Project) | [GitHub](#)

WeaveHacks 4

- Architected an **agentic** pipeline using multi-agent orchestration in **LangGraph**, with a **CEO agent** coordinating **5 sub-agents** across **OpenAI models**, incorporating **W&B Weave**, providing **observability across agent outputs and execution traces**.